

KAGGLE COMPETITION

# M5 Forecasting Accuracy

Forecasting 28 Days of Walmart Unit Sales with Machine Learning

Chingfen Hung  
Davey Johnson  
Ko-Jung Hsu  
Evelyn Lai  
Tzu-Yu Chen

**3,049**

Unique Products Tracked

**28 Days**

Forecast Horizon

**0.59212**

Best Kaggle Score

# Business Context

## THE CHALLENGE

Forecast 28 days of unit sales for Walmart product-store series across 10 stores in 3 U.S. states, using historical sales, calendar events, SNAP days, and price information.

## WHY IT MATTERS



### Inventory Efficiency

Reduce overstock and stockout risk through better demand visibility



### Proactive Planning

Help teams plan four weeks ahead instead of reacting to shortages



### Cost Reduction

Lower waste, holding costs, and unnecessary replenishment



### Promotion & Event Readiness

Capture demand shifts from price changes, SNAP days, holidays, and local events

# Data Cleaning & Preparation

INPUT → OUTPUT  
3 CSVs → 1 Clean File

## PIPELINE STEPS:

1

### Parse types & fill missing values

Dates parsed to datetime, 1,807 event NaNs filled

2

### Downcast columns

Sales int64 → int16, prices float64 → float32

3

### Detect leading zeros

77% of series have pre-shelf zeros flagged

4

### Melt wide → long & merge

Combined features into one row-per-item-day table

5

### Forward-fill price gaps

Items with late price records filled using first known price

6

### Export clean parquet

Single file for all downstream modeling notebooks

## KEY NUMBERS:

16M+

Rows After  
Melt

3,049

Unique Products  
Tracked

60%

Of Products Had  
No Daily Sales

## ISSUES FOUND & RESOLVED



**1,807 event NaNs** — not missing data, just "no event"; filled with no\_event variable



**77% leading zeros** — items not yet on shelf; masked with is\_active flag



**Price gaps after melt** — items with late price records; forward/back-filled per unique product

# Baseline Model Building - LightGBM

BASELINE RMSE

2.1619

## WHAT IS LightGBM?



LightGBM is a gradient boosting model that designed to handle large-scale tabular data fast, and natively supports categorical features without manual encoding.



Splits data into  
decision trees



Corrects errors  
step by step



Outputs sales  
prediction per row

## FEATURE LOGIC

**Categorical**      Handled natively by LightGBM

**Calendar**      Captures weekly cycles, seasonality,  
and government benefit days

**Price / Promo**      Models consumer price sensitivity and  
promotional lift

**Lag & Rolling**      Encodes recent demand momentum

## MODEL SETUP

Objective

**Tweedie**

Suited for sparse count data with many zeros

Sample

**2,000 products**

× 500 days

Features

**14 total**

Lag, rolling means, calendar, price signals

Metric

**RMSE**

Root mean squared error; lower = better

Max Rounds

**2,000**

Early stopping @ 100 rounds

Learn Rate

**0.05**

Found via Bayesian optimization

Num Leaves

**127**

Controls tree complexity

## METHODOLOGY

Iteratively add feature groups to the baseline model, then evaluate each on a held-out validation window using RMSE, with all lag and rolling features computed shift-first to prevent data leakage.

## KEY RESULT

Three extra features (snap\_flag, is\_event, lag\_56) beat complex combinations of other features.

## FINAL FEATURE TABLE

### CATEGORICAL

item\_id      Item identifier

dept\_id      Department ID

cat\_id      Category ID

store\_id      Store identifier

state\_id      State identifier

### CALENDAR

dow      Day of week

month      Month of year

snap\_flag      SNAP benefit day

is\_event      Holiday / event

### PRICE / PROMO

sell\_price      Item selling price

price\_change      Price change vs.  
prev day

is\_promo      Promotion flag

### LAG & ROLLING

lag\_7      7-day lagged  
sales

lag\_28      28-day lagged  
sales

lag\_56      56-day lagged  
sales

rmean\_7      7-day rolling  
mean

rmean\_28      28-day rolling  
mean

● New features added on top of baseline

# Fine-Tuning: Finding the Best Model Settings

LightGBM RMSE

2.0026

**Approach: 2-Phase Search** Sample: 1,000 products × 500 days

| Phase   | Method                                         | Trials    |
|---------|------------------------------------------------|-----------|
| Phase 1 | Manual Grid Search                             | 12 sets   |
| Phase 2 | Automated Smart Search (Bayesian Optimization) | 40 trials |

## Best Parameters Found

| Parameter        | Value     | Meaning                     |
|------------------|-----------|-----------------------------|
| Learning Rate    | 0.042     | How fast the model learns   |
| Tree Complexity  | 91 leaves | Level of detail captured    |
| Demand Pattern   | 1.28      | Tuned for sparse sales data |
| Min Data in Leaf | 85        | Lower = better              |

# Model Exploration: Ridge Regression

RIDGE REGRESSION RMSE

2.1180

## WHAT IS RIDGE REGRESSION?



Ridge regression is a linear model which incorporates regularization, it finds the best straight-line fit while making sure the model can still predict well with new data.



Fits a line across a graph



Shrinks noise within data



Predicts next-day sales per product

### Ridge Model Setup

- Alpha = 0.5 (regularization strength)
- StandardScaler applied to all features
- Same engineered features as team models
- Same validation window for comparison

## MODEL COMPARISON

Lower RMSE = better

### Ridge Regression

Linear baseline

2.1180

### Random Forest

Tree ensemble

2.0802

### GRU

Deep learning

2.2583

## Takeaways

### Strength

- Fast to train, fully interpretable

### Weaknesses

- Assumes linear relationships
- Produces negative predictions

# Model Exploration: Random Forest

RANDOM FOREST RMSE

2.0802

## MODEL COMPARISON

Lower RMSE = better

### Ridge Regression

Linear baseline

2.1180

### Random Forest

Tree ensemble

2.0802

### GRU

Deep learning

2.2583

## Takeaways

### Strength

- Robust to outliers and skewed data

### Weaknesses

- Slower to train than Ridge or GRU
- Cannot predict beyond the training range

## WHAT IS RANDOM FOREST?

Random Forest builds hundreds of decision trees on random subsets of data, then averages their predictions, can be seen as polling many experts instead of trusting one.



Grows 300 decision trees



Each tree sees random features



Averages votes for final prediction

## Random Forest Model Setup

- 300 trees, max depth 20
- Min 20 samples per leaf
- Same engineered features as team models
- Same validation window for comparison



# Model Exploration: GRU

GRU RMSE  
**2.2583**

## WHAT IS GRU?



A Gated Recurrent Unit (GRU) is a neural network that reads historical sales as a sequence — much like how a human analyst reviews past weeks before making a forecast.



Reads 28-day  
history window



Learns demand  
patterns over time



Predicts next-day  
sales per product

## GRU Model Setup

- 28-day input window
- Same engineered features as team models
- Same validation window for comparison

## MODEL COMPARISON

Lower RMSE = better

### Ridge Regression

Linear baseline

2.1180

### Random Forest

Tree ensemble

2.0802

### GRU

Deep learning

**2.2583**

## Takeaways

### Strength

- Captures sequential demand patterns

### Weaknesses

- Less effective on days with zero sales
- Performed the worst out of all models

# Final Prediction: LightGBM

KAGGLE SCORE

0.59212

## Recursive Forecasting:

1. Get static features (calendar, price, events)
2. Compute lag\_7, lag\_28, lag\_56 from history
3. Compute rmean\_7, rmean\_28 from history
4. Predict demand for each ID
5. Clip negative predictions to 0

Append predictions to history → repeat


💡 Enables multi-step prediction without retraining the model

## Ensemble Strategy:

**Train multiple models with three different seeds**

👉 Reduces variance & Improves stability

📈 ~0.001 WRMSSE improvement vs single model



The model learns both temporal dynamics and item-level behavior

# Recommendations

## Business Impact

- Improves visibility into future demand of 3,049 products sold across 10 stores
- Helps inventory teams plan and allocate stock earlier
- Reduces waste from over- or under-preparation

## Risks & Limitations

- Sudden demand shifts (e.g. supply disruptions) may reduce accuracy
- Promotions or local events may need manual adjustment
- Forecast uncertainty increases further out in the 28-day horizon

## Next Steps

- Refresh the model periodically with new sales data
- Incorporate additional business context (e.g. pricing, events) into future forecasts
- Use forecast output alongside manager judgment — not as a replacement